

# C++ - parallelization and synchronization

---

Jakub Yaghob

Martin Kruliš





# The problem

- Race conditions
  - Separate threads with shared state
  - Result of computation depends on OS scheduling

# Race conditions – simple demo



- Linked list
- Shared state

```
List lst;
```


- Thread A

```
lst.push_front(A);
```


- Thread B

```
lst.push_front(B);
```


Initial state




Correct state



Another correct state



Incorrect state



# Race conditions – advanced demo



```
struct Counter {  
    Counter():value(0) { }  
    int value;  
    void increment()  
    { ++value; }  
    void decrement()  
    { --value; }  
    int get()  
    { return value; }  
};
```

- Shared state  
`Counter c;`
- Thread A  
`c.increment();`  
`cout << c.get();`
- Thread B  
`c.increment();`  
`cout << c.get();`
- Possible outputs  
12, 21, **11**



# C++ 11 features

- Atomic operations
- Low-level threads
- High-level futures
- Synchronization primitives
- Thread-local storage



# C++ 14 and C++17 features

- C++14 features
  - Shared timed mutex
- C++17 features
  - Parallel algorithms
  - Shared mutex



# C++ 11 – atomic operations

- Atomic operations
  - Header `<atomics>`
  - Allows creating portable lock-free algorithms and data structures
  - Memory ordering
  - Fences
  - Lock-free operations, algorithms, data-structures



# C++ 11 – atomic operations

- Memory ordering

- `enum memory_order;`

- `memory_order_seq_cst`

- Sequentially consistent, most restrictive memory model

- `memory_order_relaxed`

- Totally relaxed memory model, allows best freedom for CPU and compiler optimizations

- `memory_order_acquire`, `memory_order_release`,  
`memory_order_acq_rel`

- Additional barriers, weaker than sequentially consistent, stronger than relaxed





# C++ 11 – atomic operations

- Barriers
  - Acquire barrier
    - All loads read after acquire will perform after it (loads do not overtake acquire)
  - Release barrier
    - All stores written before release are committed before the release (writes do not delay)



# C++ 11 – atomic operations

- Easy way to make the demo safe

```
#include <atomic>
```

```
struct Counter {  
    std::atomic<int> value;  
    void increment(){ ++value; }  
    void decrement(){ --value; }  
    int get(){ return value.load(); }  
};
```



# C++ 11 – atomic operations

- Template atomic
  - Defined for any type
    - Load, store, compare\_exchange
  - Specialized for bool, all integral types, and pointers
    - Load, store, compare\_exchange
    - Arithmetic and bitwise operations
      - fetch\_add



# C++ 11 – atomic operations

- Atomic flag
  - `atomic_flag` allows one-bit test and set
- Atomic operations for `shared_ptr`

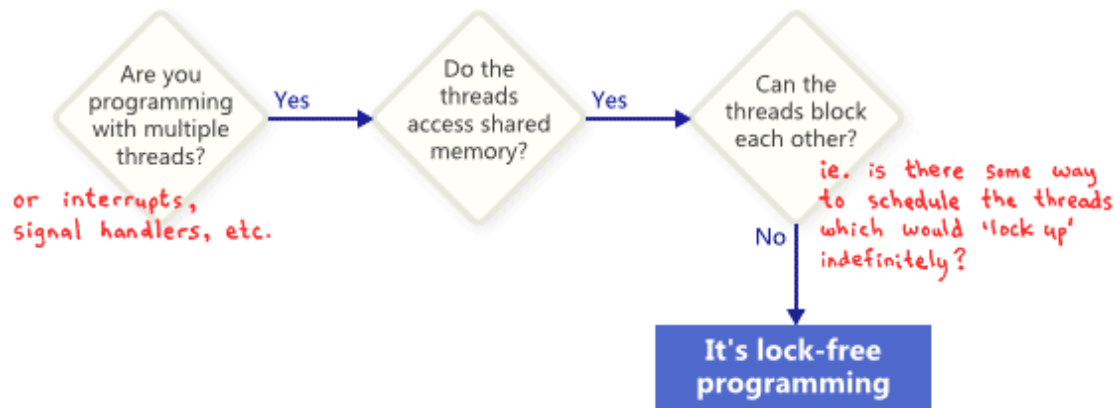


# C++ 11 – atomic operations

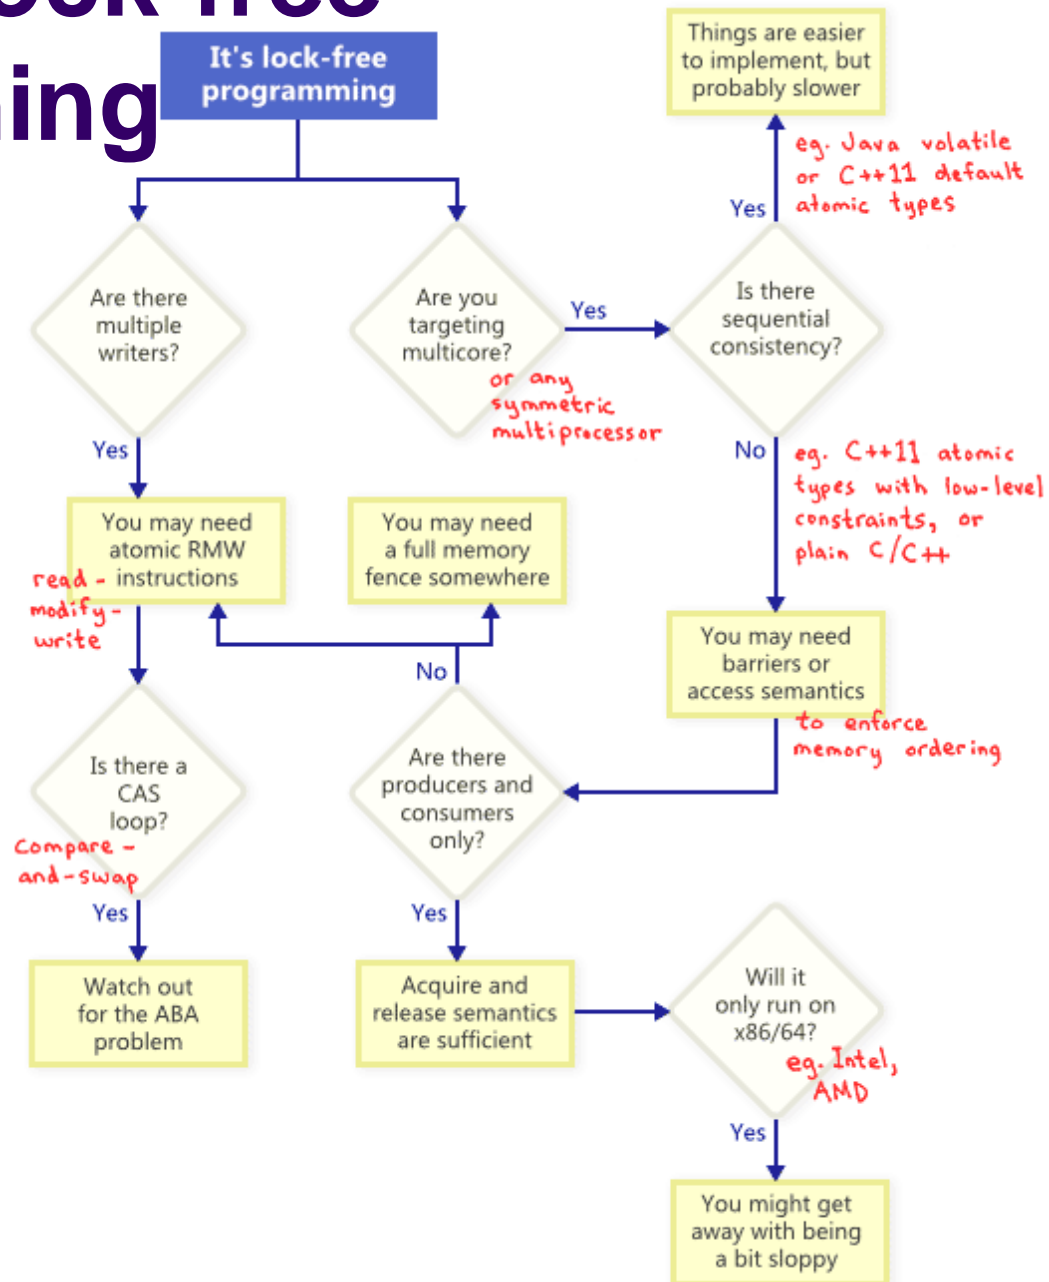
- Fences

- Explicit memory barrier
- `void atomic_thread_fence(memory_order order) noexcept;`
  - `memory_order_relaxed`
    - No effect
  - `memory_order_acquire`
    - An acquire fence
  - `memory_order_release`
    - A release fence
  - `memory_order_acq_rel`
    - Both an acquire and a release fence
  - `memory_order_seq_cst`
    - Sequentially consistent

# C++ 11 – lock-free programming



# C++ 11 – lock-free programming





# C++ 11 – threads

- Low-level threads
  - Header `<thread>`
  - `thread` class
  - Fork-join paradigm
  - Namespace `this_thread`





# C++ 11 – threads

- Class thread
  - Constructor
    - `template <class F, class ...Args>  
explicit thread(F&& f, Args&&... args);`
  - Destructor
    - If `joinable()` then `terminate()`
  - `bool joinable() const noexcept;`
  - `void join();`
    - Blocks, until the thread `*this` has completed
  - `void detach();`
  - `id get_id() const noexcept;`
  - `static unsigned hardware_concurrency();`



# C++ 11 – threads

- Namespace `this_thread`
  - `thread::id get_id() noexcept;`
    - Unique ID of the current thread
  - `void yield() noexcept;`
    - Opportunity to reschedule
  - `sleep_for, sleep_until`
    - Blocks the thread for relative/absolute timeout



# C++ 11 – threads

- Demo

```
#include <iostream>
```


```
#include <thread>
```

```
void thread_fn() { std::cout << "Hello from thread" <<  
    std::endl; }
```

```
int main(int argc, char **argv) {  
    std::thread thr(&thread_fn);  
    std::cout << "Hello from main" << std::endl;  
    thr.join();  
    return 0;  
}
```

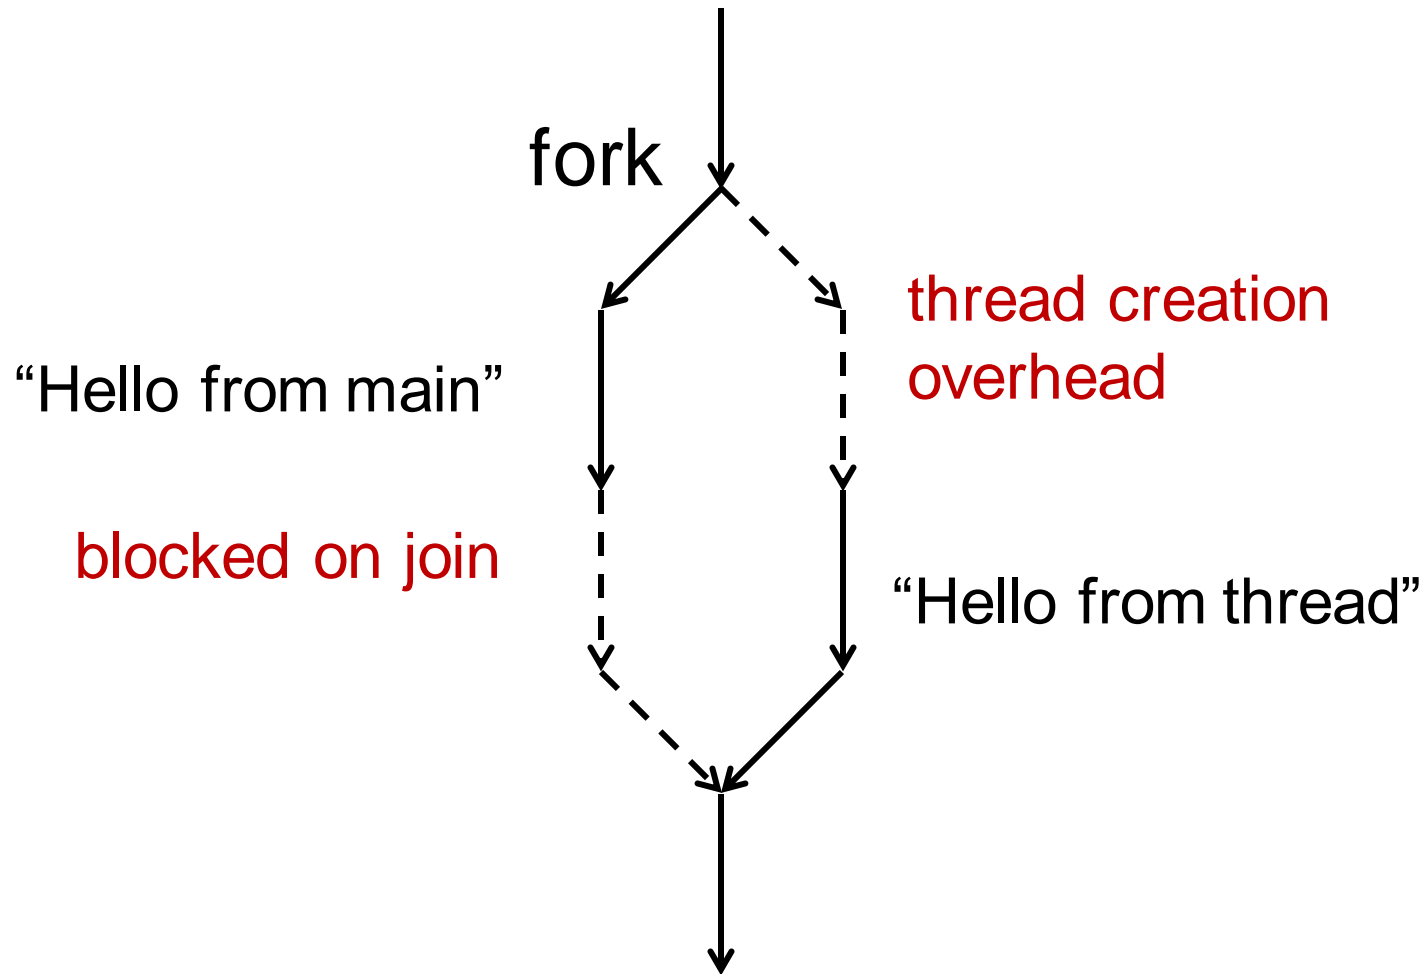


# C++ 11 – threads






# C++ 11 – threads



# C++ 11 – threads





# C++ 11 – threads

## ● Demo

```
#include <iostream>
#include <thread>
#include <vector>

int main(int argc, char **argv) {
    std::vector<std::thread> workers;
    for(int i=0;i<10;++i)
        workers.push_back(std::thread([i]() {
            std::cout << "Hello from thread " << i << std::endl;
        }));
    std::cout << "Hello from main" << std::endl;
    for(auto &t : workers)
        t.join();
    return 0;
}
```



# C++ 11 – threads

- Passing arguments to threads
  - By value
    - Safe, but you MUST make deep copy
  - By move (rvalue reference)
    - Safe, as long as strict (deep) adherence to move semantics
  - By const reference
    - Safe, as long as object is guaranteed deep-immutable
  - By non-const reference
    - Safe, as long as the object is monitor





# C++ 11 – futures

- Futures
  - Header `<future>`
  - High-level asynchronous execution
  - Future
  - Promise
  - Async
  - Error handling



# C++ 11 – futures

- Shared state
  - Consist of
    - Some state information and some (possibly not yet evaluated) result, which can be a (possibly void) value or an exception
  - Asynchronous return object
    - Object, that reads results from an shared state
  - Waiting function
    - Potentially blocks to wait for the shared state to be made ready
  - Asynchronous provider
    - Object that provides a result to a shared state



# C++ 11 – futures

- Future
  - `std::future<T>`
  - Future value of type T
  - Retrieve value via `get()`
    - Waits until the shared state is ready
  - `wait()`, `wait_for()`, `wait_until()`
  - `std::shared_future<T>`
    - Value can be read by more than one thread



# C++ 11 – futures

- Async
  - `std::async`
  - Higher-level convenience utility
  - Launches a function potentially in a new thread

- Async usage

```
int foo(double, char, bool);  
auto fut = std::async(foo, 1.5, 'x', false);  
auto res = fut.get();
```



# C++ 11 – futures

- Packaged task
  - `std::packaged_task`
  - How to implement async with more control
  - Wraps a function and provides a future for the function result value, but the object itself is callable



# C++ 11 – futures

- Packaged task usage

```
std::packaged_task<int(double, char, bool)>  
    tsk(foo);  
auto fut = tsk.get_future();  
std::thread thr(std::move(tsk), 1.5, 'x', false);  
auto res = fut.get();
```



# C++ 11 – futures

- Promise

- `std::promise<T>`

- Lowest-level

- Steps

- Calling thread makes a promise
    - Calling thread obtains a future from the promise
    - The promise, along with function arguments, are moved into a separate thread
    - The new thread executes the function and fulfills the promise
    - The original thread retrieves the result



# C++ 11 – futures

- Promise usage

- Thread A

```
std::promise<int> prm;  
auto fut = prm.get_future();  
std::thread thr(thr_fnc, std::move(prm));  
auto res = fut.get();
```

- Thread B

```
void thr_fnc(std::promise<int> &&prm) {  
    prm.set_value(123);  
}
```





# C++ 11 – futures

- Constraints

- A default-constructed promise is inactive
  - Can die without consequence
- A promise becomes active, when a future is obtained via `get_future()`
  - Only one future may be obtained
- A promise must either be satisfied via `set_value()`, or have an exception set via `set_exception()`
  - A satisfied promise can die without consequence
  - `get()` becomes available on the future
  - A promise with an exception will raise the stored exception upon call of `get()` on the future
  - A promise with neither value nor exception will raise “broken promise” exception



# C++ 11 – futures

- Exceptions

- All exceptions of type `std::future_error`
  - Has error code with enum type `std::future_errc`

- inactive promise

```
std::promise<int> pr;  
// fine, no problem
```

- active promise, unused

```
std::promise<int> pr;  
auto fut = pr.get_future();  
// fine, no problem  
// fut.get() blocks indefinitely
```

- too many futures

```
std::promise<int> pr;  
auto fut1 = pr.get_future();  
auto fut2 = pr.get_future();  
// error "Future already  
retrieved"
```



# C++ 11 – futures

- satisfied promise

```
std::promise<int> pr;  
auto fut = pr.get_future();  
{ std::promise<int>  
pr2(std::move(pr));  
    pr2.set_value(10);  
}  
auto r = fut.get();  
// fine, return 10
```

- too much satisfaction

```
std::promise<int> pr;  
auto fut = pr.get_future();  
{ std::promise<int>  
pr2(std::move(pr));  
    pr2.set_value(10);  
    pr2.set_value(11);  
// error "Promise already  
satisfied"  
}  
auto r = fut.get();
```



# C++ 11 – futures

- exception

```
std::promise<int> pr;  
auto fut = pr.get_future();  
{ std::promise<int> pr2(std::move(pr));  
  pr2.set_exception(  
    std::make_exception_ptr(  
      std::runtime_error("bububu"))) ;  
}  
auto r = fut.get();  
// throws the runtime_error
```



# C++ 11 – futures

- broken promise

```
std::promise<int> pr;  
auto fut = pr.get_future();  
{ std::promise<int> pr2(std::move(pr));  
  // error "Broken promise"  
}  
auto r = fut.get();
```

# C++ 11 – synchronization primitives



- Synchronization primitives
  - Mutual exclusion
    - Header <mutex>
  - Condition variables
    - Header <condition\_variable>



# C++ 11 – mutex

- Mutex
  - A synchronization primitive that can be used to protect shared data from being simultaneously accessed by multiple threads
  - `mutex` offers exclusive, non-recursive ownership semantics
    - A calling thread owns a `mutex` from the time that it successfully calls either `lock` or `try_lock` until it calls `unlock`
    - When a thread owns a `mutex`, all other threads will block (for calls to `lock`) or receive a false return value (for `try_lock`) if they attempt to claim ownership of the `mutex`
    - A calling thread must not own the `mutex` prior to calling `lock` or `try_lock`
  - The behavior of a program is undefined if a `mutex` is destroyed while still owned by some thread



# C++ 11 – mutex example

- Shared state

```
List lst;  
std::mutex mtx;
```

- Thread A

```
mtx.lock();  
lst.push_front(A);  
mtx.unlock();
```

- Thread B

```
mtx.lock();  
lst.push_front(B);  
mtx.unlock();
```





# C++ 11 – mutex variants

- Other `mutex` variants

- `timed_mutex`

- In addition, `timed_mutex` provides the ability to attempt to claim ownership of a `timed_mutex` with a timeout via the `try_lock_for` and `try_lock_until`

- `recursive_mutex`

- exclusive, recursive ownership semantics
      - A calling thread *owns* a `recursive_mutex` for a period of time that starts when it successfully calls either `lock` or `try_lock`. During this period, the thread may make additional calls to `lock` or `try_lock`. The period of ownership ends when the thread makes a matching number of calls to `unlock`
      - When a thread owns a `recursive_mutex`, all other threads will block (for calls to `lock`) or receive a false return value (for `try_lock`) if they attempt to claim ownership of the `recursive_mutex`
      - The maximum number of times that a `recursive_mutex` may be locked is unspecified, but after that number is reached, calls to `lock` will throw `std::system_error` and calls to `try_lock` will return false

- `recursive_timed_mutex`

- Combination



# C++ 11 – mutex wrappers

- `std::unique_lock`
  - Lock class with more features
    - Timed wait, deferred lock
- `std::lock_guard`
  - Scope based lock (RAII)
  - Linked list demo, code for one thread

```
{  
std::lock_guard<std::mutex> lk(mtx);  
lst.push_front(X);  
}
```

# C++ 14 – mutex variants and wrappers



- Other mutex variants in C++ 14
  - `std::shared_timed_mutex`
    - Multiple threads can make shared lock using `lock_shared()`
- Additional wrapper
  - `std::shared_lock`
    - Calls `lock_shared` for the given mutex

# C++ 17 – mutex variants, wrappers, and others



- Another mutex variant
  - `std::shared_mutex`
- Variadic wrapper
  - `template <typename ... MutexTypes> class scoped_lock;`
    - Multiple locks at once
- Interference size
  - `std::size_t hardware_destructive_interference_size;`
  - Size of a cache line



# C++ 11 – locking algorithms

- `std::lock`
  - locks specified mutexes, blocks if any are unavailable
- `std::try_lock`
  - attempts to obtain ownership of mutexes via repeated calls to `try_lock`

```
// don't actually take the locks yet
std::unique_lock<std::mutex> lock1(mtx1, std::defer_lock);
std::unique_lock<std::mutex> lock2(mtx2, std::defer_lock);
// lock both unique_locks without deadlock
std::lock(lock1, lock2);
```



# C++ 11 – call once

- `std::once_flag`
  - Helper object for `std::call_once`
- `std::call_once`
  - invokes a function only once even if called from multiple threads

```
std::once_flag flag;  
void do_once() {  
    std::call_once(flag, []() { do something only once });  
}  
std::thread t1(do_once);  
std::thread t2(do_once);
```



# C++ 11 – condition variable

- `std::condition_variable`
  - Can be used to block a thread, or multiple threads at the same time, until
    - a notification is received from another thread
    - a timeout expires, or
    - a spurious wakeup occurs
      - Appears to be signaled, although the condition is not valid
      - Verify the condition after the thread has finished waiting
  - Works with `std::unique_lock`
  - `wait` atomically manipulates mutex, `notify` does nothing

# C++11 – condition variable example



```
std::mutex m;  
std::condition_variable cond_var;  
bool done = false; bool notified = false;
```

## ● Producer

```
for () {  
    // produce something  
    { std::lock_guard<std::mutex>  
        lock(m);  
        queue.push(item);  
        notified = true; }  
    cond_var.notify_one();  
}  
std::lock_guard<std::mutex> lock(m);  
notified = true;  
done = true;  
cond_var.notify_one();
```

## ● Consumer

```
std::unique_lock<std::mutex> lock(m);  
while(!done) {  
    while (!notified) {  
        // loop to avoid spurious wakeups  
        cond_var.wait(lock);  
    }  
    while(!produced_nums.empty()) {  
        // consume  
        produced_nums.pop();  
    }  
    notified = false;  
}
```





# C++ 11 – thread-local storage

- Thread-local storage
  - Added a new storage-class
  - Use keyword **thread\_local**
    - Must be present in all declarations of a variable
    - Only for namespace or block scope variables and to the names of static data members
      - For block scope variables **static** is implied
  - Storage of a variable lasts for the duration of a thread in which it is created



# C++ extensions – parallelism

- Parallelism
  - TS v1 adopted in C++ 17, TS v2 finished
    - In headers `<algorithm>`, `<numeric>`
  - Parallel algorithms
    - Execution policy in `<execution>`
      - `seq` – execute sequentially
      - `par` – execute in parallel on multiple threads
      - `par_unseq` – execute in parallel on multiple threads, interleave individual iterations within a single thread, no locks
      - `unseq` – (C++20) execute in single thread+vectorized
    - `for_each`
    - `reduce`, `scan`, `transform_reduce`, `transform_scan`
      - Inclusive scan – like `partial_sum`, includes i-th input element in the i-th sum
      - Exclusive scan – like `partial_sum`, excludes i-th input element from the i-th sum
    - No exceptions should be thrown
      - Terminate

# C++ extensions – parallelism

## v1



- Parallel algorithms
  - Not all algorithms have parallel version
  - `adjacent_difference`, `adjacent_find`, `all_of`, `any_of`, `copy`, `copy_if`, `copy_n`, `count`, `count_if`, `equal`, `exclusive_scan`, `fill`, `fill_n`, `find`, `find_end`, `find_first_of`, `find_if`, `find_if_not`, `for_each`, `for_each_n`, `generate`, `generate_n`, `includes`, `inclusive_scan`, `inner_product`, `inplace_merge`, `is_heap`, `is_heap_until`, `is_partitioned`, `is_sorted`, `is_sorted_until`, `lexicographical_compare`, `max_element`, `merge`, `min_element`, `minmax_element`, `mismatch`, `move`, `none_of`, `nth_element`, `partial_sort`, `partial_sort_copy`, `partition`, `partition_copy`, `reduce`, `remove`, `remove_copy`, `remove_copy_if`, `remove_if`, `replace`, `replace_copy`, `replace_copy_if`, `replace_if`, `reverse`, `reverse_copy`, `rotate`, `rotate_copy`, `search`, `search_n`, `set_difference`, `set_intersection`, `set_symmetric_difference`, `set_union`, `sort`, `stable_partition`, `stable_sort`, `swap_ranges`, `transform`, `transform_exclusive_scan`, `transform_inclusive_scan`, `transform_reduce`, `uninitialized_copy`, `uninitialized_copy_n`, `uninitialized_fill`, `uninitialized_fill_n`, `unique`, `unique_copy`

# C++ extensions – parallelism v2



- Task block
  - Support for fork-join paradigm
  - Spawn other task\_blocks and wait for their completion
  - Exceptions
    - Each task\_block has an exception list
    - Exceptions from forked task\_blocks are stored in the exception list
    - Exceptions are invoked when task\_block finishes



# C++ extension – executors

- Executors
  - Now separate TS, maybe finished in C++23 timeframe
- Executor
  - Controls how a task (=function) is executed
  - Direction
    - One-way execution
      - Does not return a result
    - Two-way execution
      - Returns future
    - Then
      - Execution agent begins execution after a given future becomes ready, returns future
  - Cardinality
    - Single
      - One execution agent
    - Bulk executions
      - Group of execution agents
      - Agents return a factory
- Thread pool
  - Controls where the task is executed



# C++ extensions – concurrency

- Concurrency
  - TS published, depends on executors TS
  - Improvements to future
    - `future<T2> then(F &&f)`
      - Execute asynchronously a function `f` when the future is ready
  - Latches
    - Thread coordination mechanism
      - Block one or more threads until an operation is completed
    - Single use
  - Barriers
    - Thread coordination mechanism
    - Reusable
    - Multiple barrier types
      - `barrier`
      - `flex_barrier` – calls a function in a completion phase

# C++ extension – transactional memory



- TS v1 finished
- Transactional memory
  - Added several keywords for statements and declarations
  - **synchronized** *compound-statement*
    - Synchronized with other synchronized blocks
    - One global lock for all synchronized blocks
  - Atomic blocks
    - Execute atomically and not concurrently with synchronized blocks
    - Can execute concurrently with other atomic blocks if no conflicts
    - Differs in behavior with exceptions
  - **atomic\_noexcept** *compound-statement*
    - Escaping exception causes undefined behavior
  - **atomic\_cancel** *compound-statement*
    - Escaping exception rolls back the transaction, but must be transaction safe
    - Functions can be declared **transaction\_safe**
  - **atomic\_commit** *compound-statement*
    - Escaping exception commits the transaction