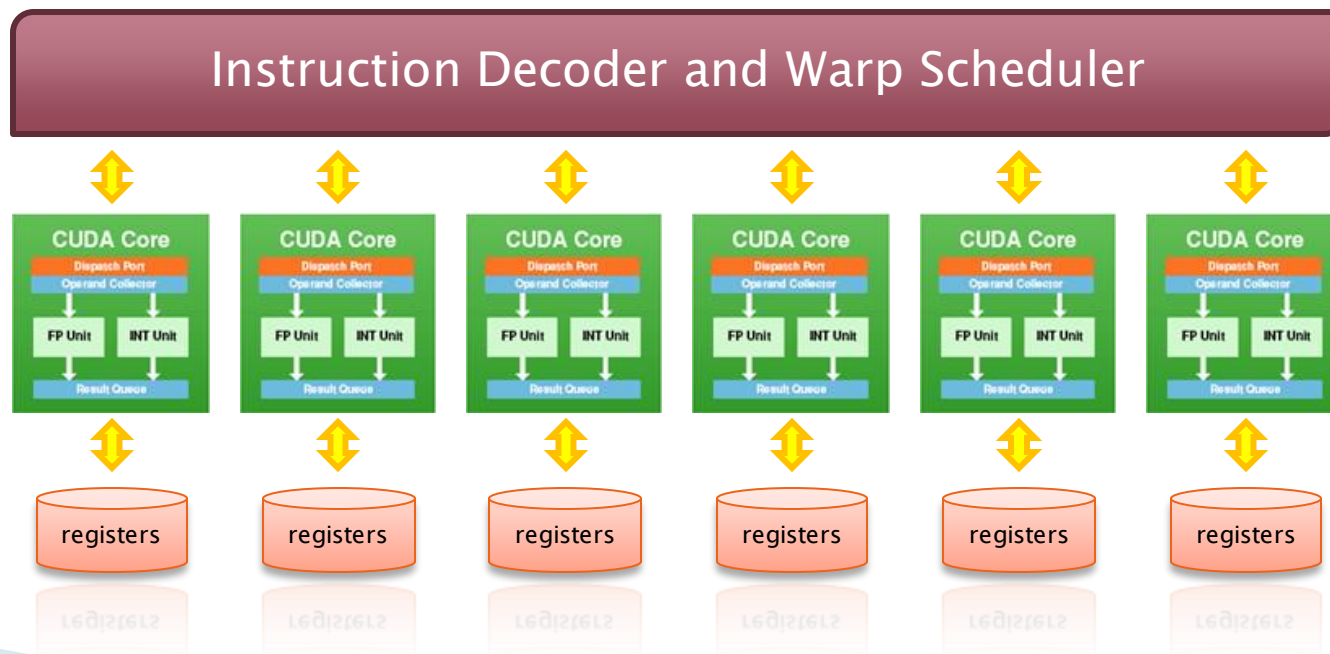


GPU Architectures and CUDA in More Detail

Martin Kruliš

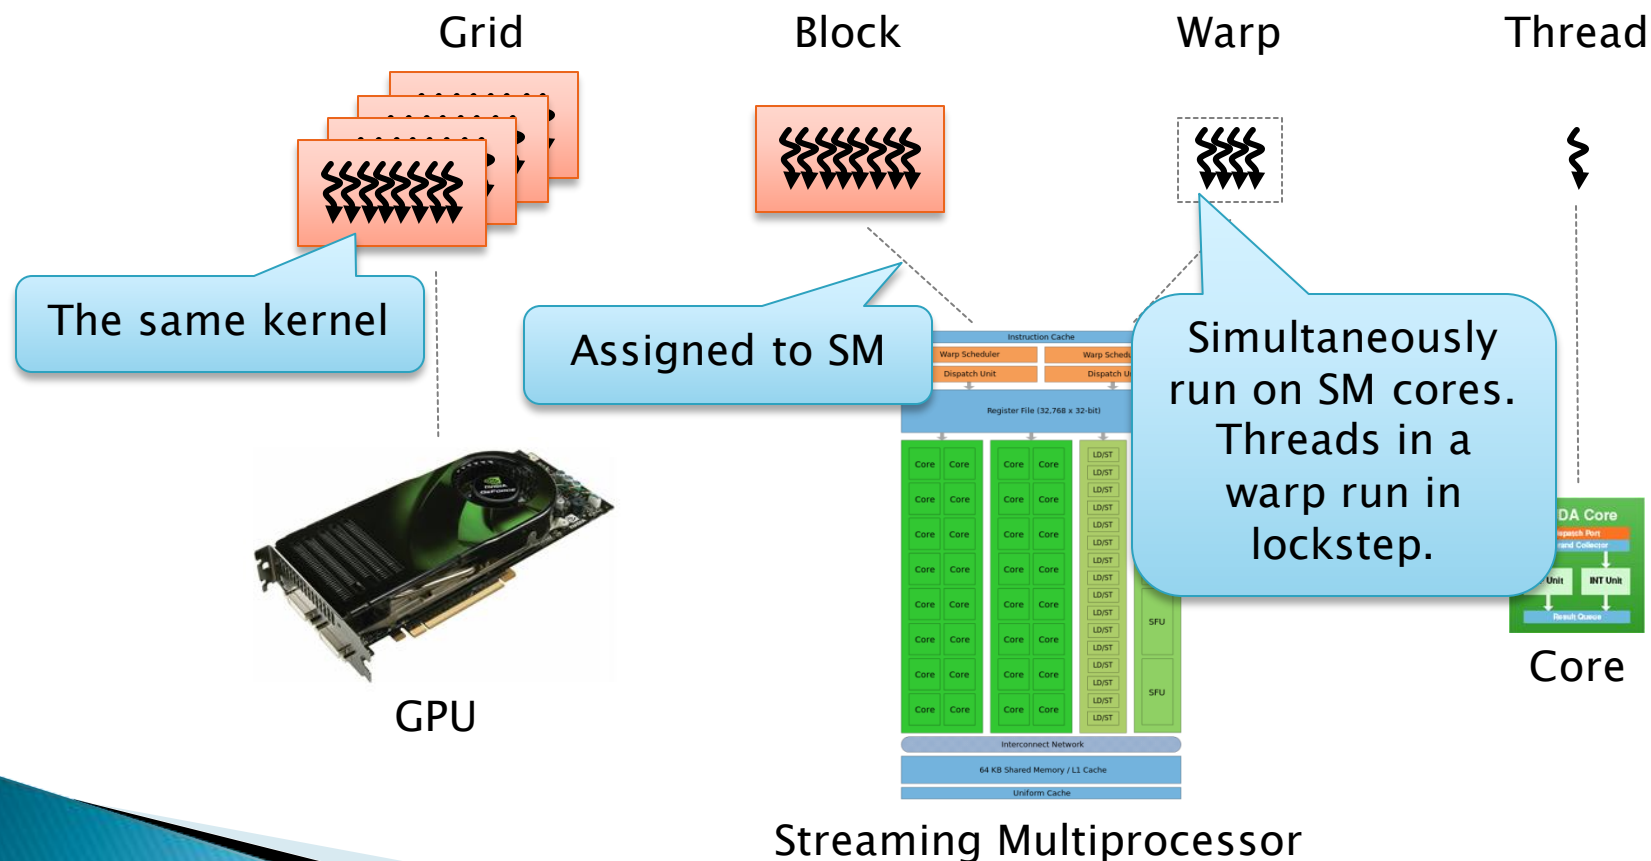
SIMT Execution (Revision)

- ▶ Single Instruction Multiple Threads
 - All cores are executing the same instruction
 - Each core has its own set of registers



Thread-Core Mapping (Revision)

- ▶ How are threads assigned to SMPs



Instruction Schedulers

► Decomposition

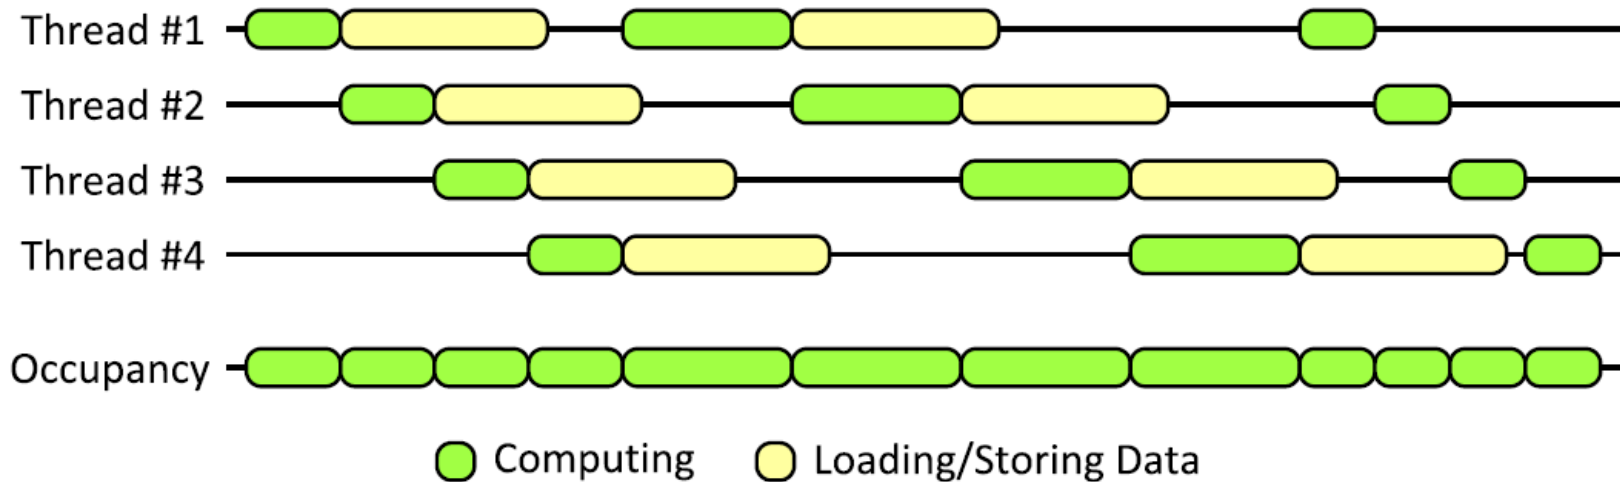
- Each block assigned to the SMP is divided into warps and the warps are assigned to schedulers

► Schedulers

- Select warp that is ready at every instruction cycle
- The SMP instruction throughput depends on CC:
 - 1.x – 1 instruction per 4 cycles, 1 scheduler
 - 2.0 – 1 instruction per 2 cycles, 2 schedulers
 - 2.1 – 2 instructions per 2 cycles, 2 schedulers
 - 3.x and 5.x – 2 instructions per cycle, 4 schedulers

Hiding Latency

- ▶ Fast Context Switch
 - When a warp gets stalled
 - E.g., by data load/store
 - Scheduler switch to next active warp




SIMT and Branches

▶ Masking Instructions

- In case of data-driven branches
 - if-else conditions, while loops, ...
- All branches are traversed, threads mask their execution in invalid branches

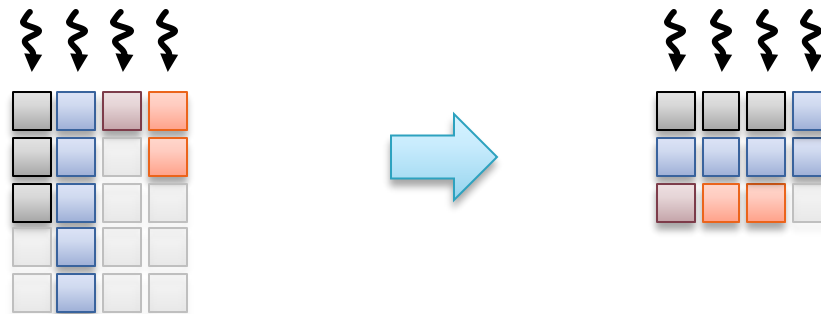
```
if (threadIdx.x % 2 == 0) {  
    ... even threads code ...  
} else {  
    ... odd threads code ...  
}
```



Reducing Thread Divergence

▶ Work Reorganization

- In case the workload is imbalanced
- Cheap balancing can lead to better occupancy



◦ Example

- Matrix with dimensions not divisible by warp size
- Item (i, j) has linear index $i * \text{width} + j$

Block-wise Synchronization

▶ Memory Fences

```
__threadfence();  
__threadfence_block();  
__threadfence_system();
```

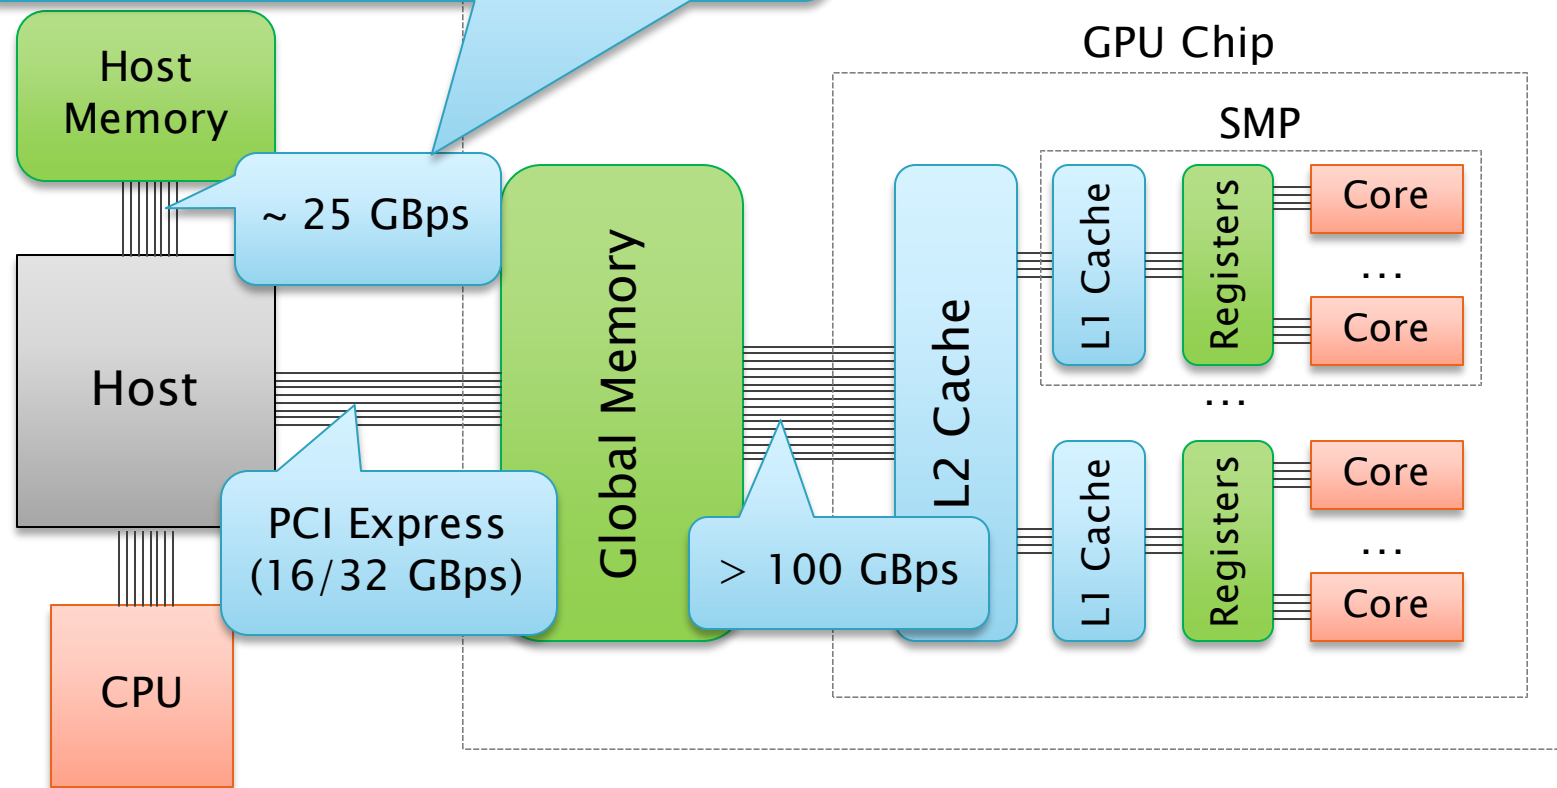
▶ Barrier

- Synchronization between warps in block

```
__syncthreads();  
__syncthreads_count(predicate);  
__syncthreads_and(predicate);  
__syncthreads_or(predicate);
```

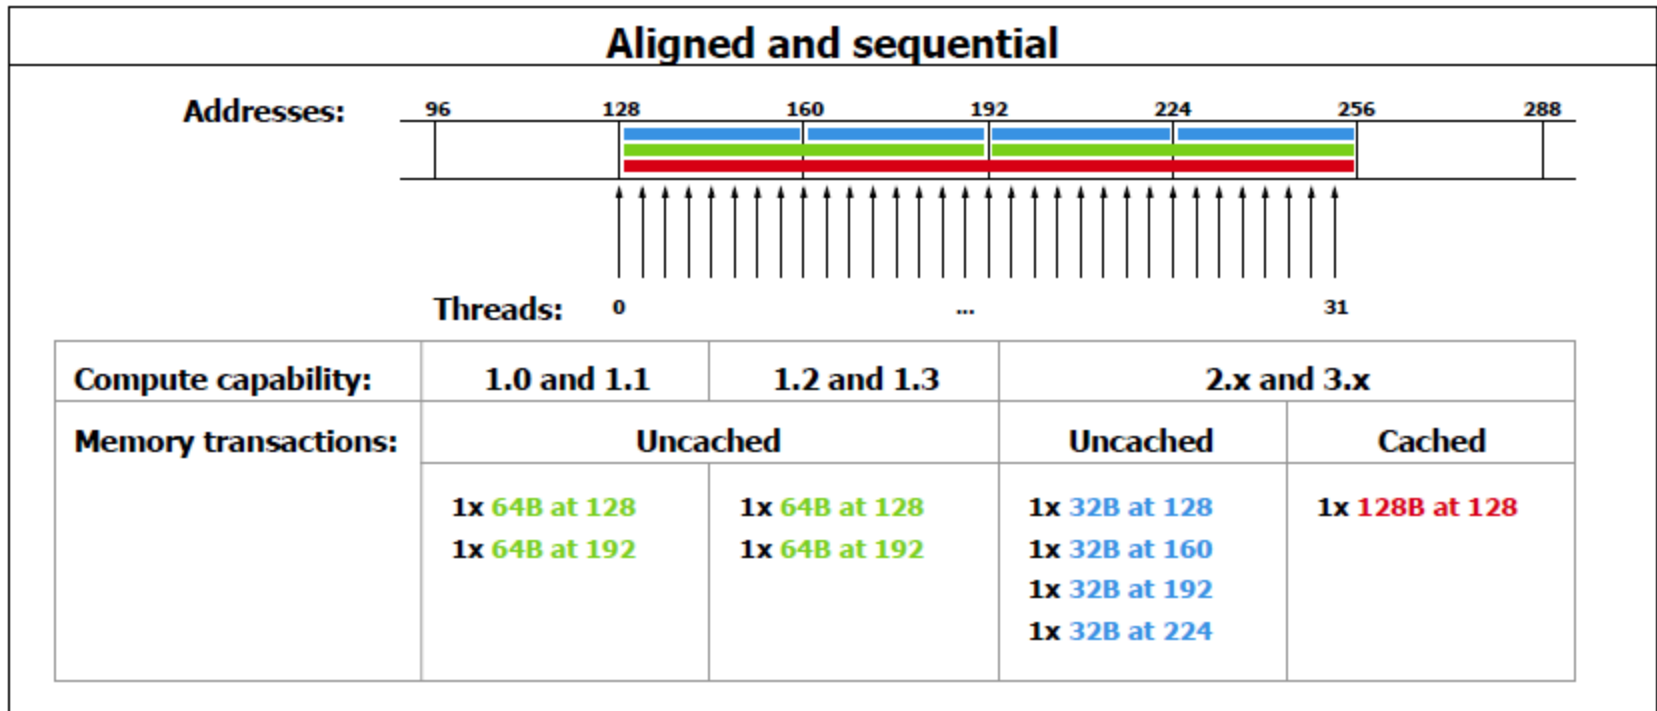

GPU Memory (Revision)

Note that details about host memory interconnection are platform specific



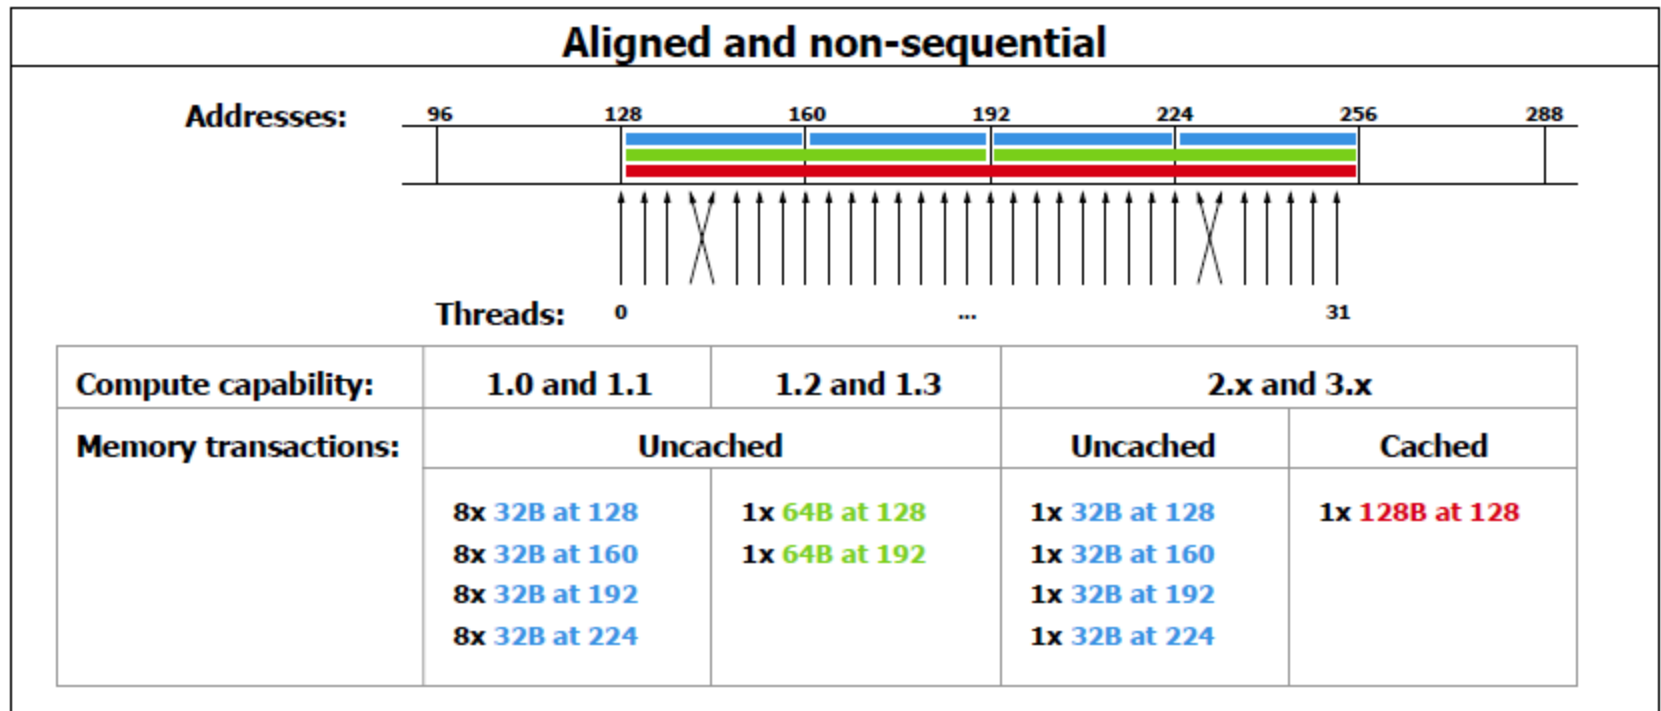
Global Memory

- ▶ Access Patterns
 - Perfectly aligned sequential access



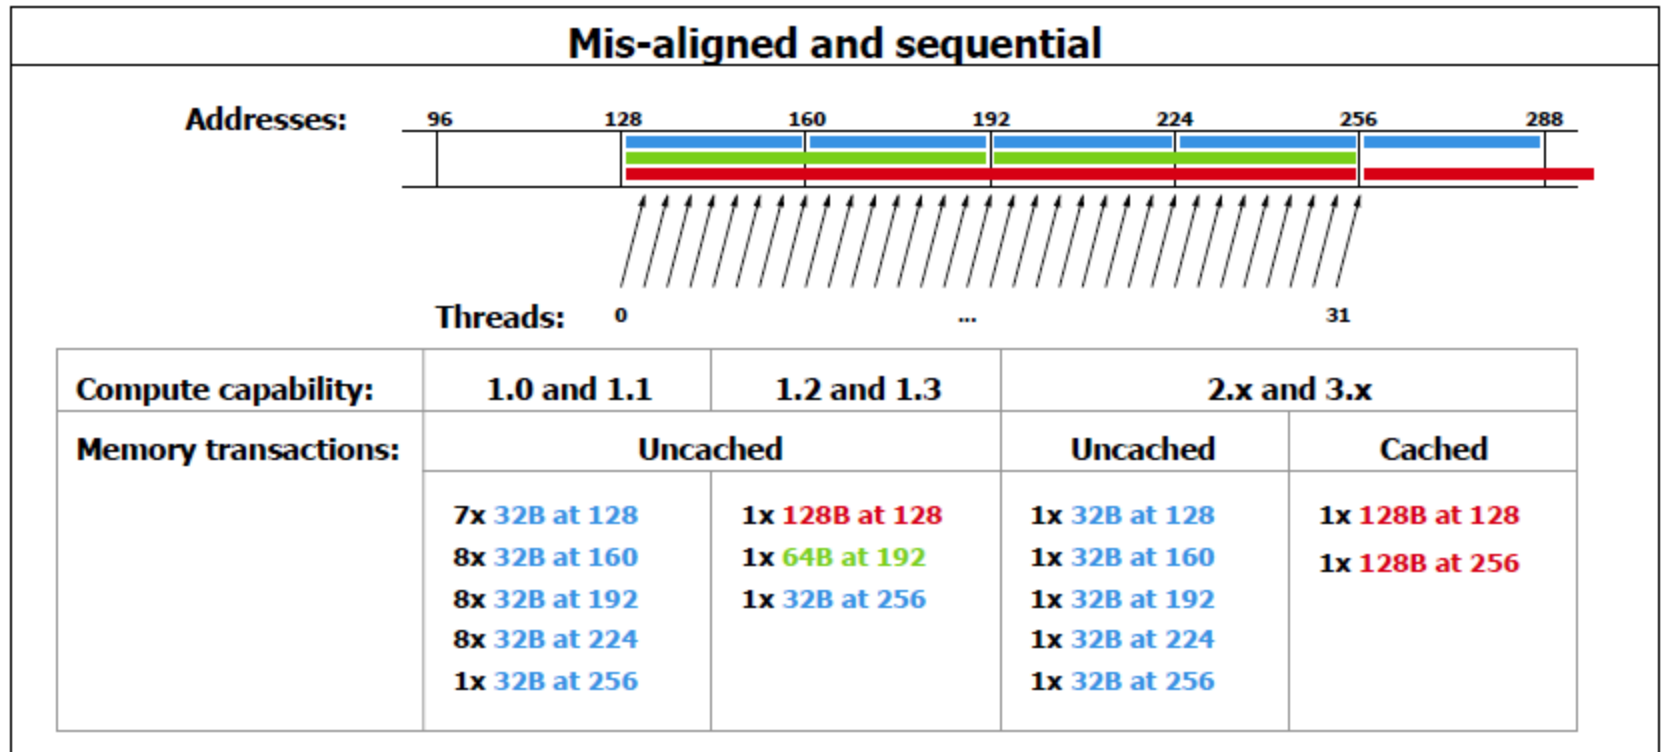
Global Memory

- ▶ Access Patterns
 - Perfectly aligned with permutation



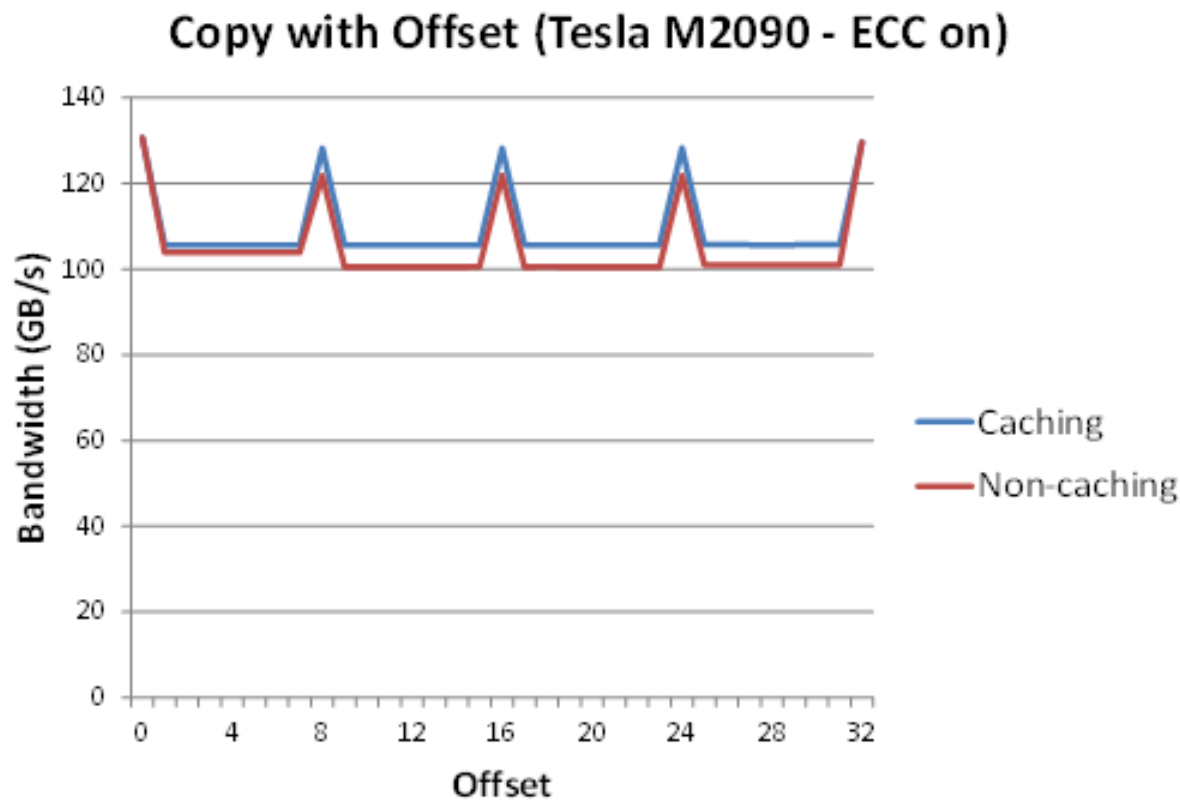
Global Memory

- ▶ Access Patterns
 - Continuous sequential, but misaligned



Global Memory

- ▶ Coalesced Loads Impact



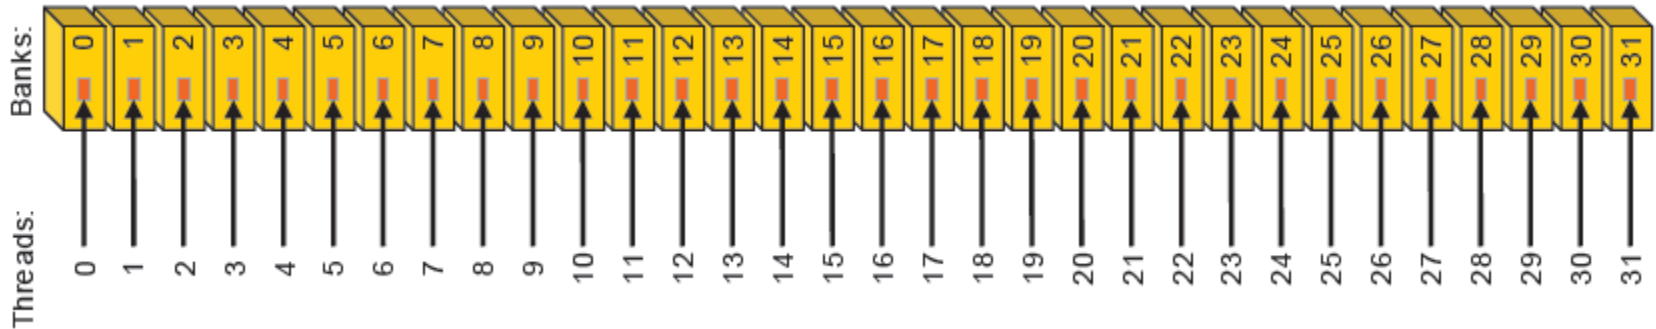
Shared Memory

- ▶ Memory Shared by SM
 - Divided into banks
 - Each bank can be accessed independently
 - Consecutive 32-bit words are in consecutive banks
 - Optionally, 64-bit words division is used (CC 3.x)
 - Bank conflicts are serialized
 - Except for reading the same address (broadcast)

Compute capability	Mem. size	# of banks	latency
1.x	16 kB	16	32 bits / 2 cycles
2.x	48 kB	32	32 bits / 2 cycles
3.x	48 kB	32	64 bits / 1 cycle

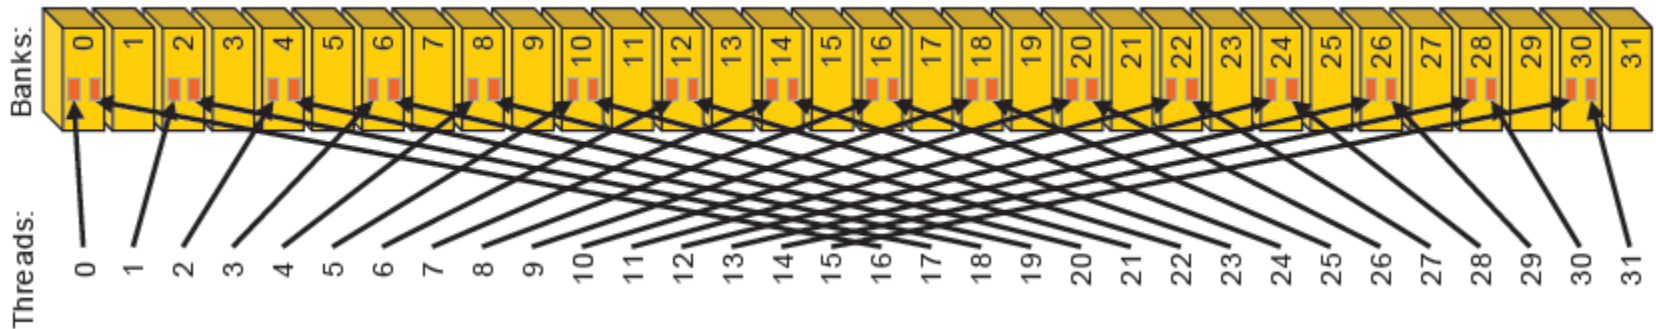
Shared Memory

- ▶ Linear Addressing
 - Each thread in warp access different memory bank
 - No collisions




Shared Memory

- ▶ Linear Addressing with Stride
 - Each thread access $2*i$ -th item
 - 2-way conflicts (2x slowdown) on CC < 3.0
 - No collisions on CC 3.x
 - Due to 64-bits per cycle throughput



Shared Memory

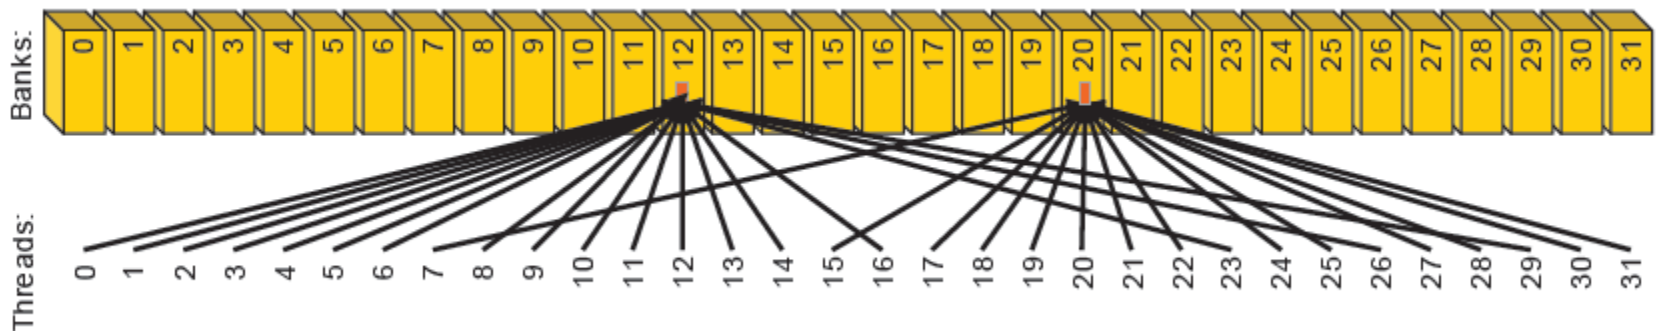
- ▶ Linear Addressing with Stride
 - Each thread access $3*i$ -th item
 - No collisions, since the number of banks is not divisible by the stride



Shared Memory

► Broadcast

- One set of threads access value in bank #12 and the remaining threads access value in bank #20
- Broadcasts are served independently on CC 1.x
 - I.e., sample below causes 2-way conflict
- CC 2.x and newer serve broadcasts simultaneously



Registers

▶ Registers

- One register pool per multiprocessor
 - 8–64k of 32-bit registers (depending on CC)
 - Register allocation is defined by compiler
- All allocated blocks share the registry pool
 - Register pressure (heavy utilization) may limit number of blocks running simultaneously
 - It may also cause registry spilling
- As fast as the cores (no extra clock cycles)
- Read-after-write dependency
 - 24 clock cycles
 - Can be hidden if there are enough active warps

Local Memory

- ▶ Per-thread Global Memory
 - Allocated automatically by compiler
 - Compiler may report the amount of allocated local memory (use `--ptxas-options=-v`)
 - Large local structures and arrays are places here
 - Instead of the registers
 - Register Pressure
 - The registers are spilled into the local memory

Memory Allocation

▶ Global Memory

- `cudaMalloc()`, `cudaFree()`
- Dynamic kernel allocation
 - `malloc()` and `free()` called from kernel
 - `cudaDeviceSetLimit(cudaLimitMallocHeapSize, size)`

▶ Shared Memory

- Statically (e.g., `__shared__ int foo[16];`)
- Dynamically (by 3rd kernel launch parameter)

```
extern __shared__ float bar[];  
float *bar1 = &(bar[0]);  
float *bar2 = &(bar[size_of_bar1]);
```

Page-locked Memory

▶ Page-locked (Pinned) Host Memory

- Host memory that is prevented from swapping
- Created/dismissed by

`cudaHostAlloc()`, `cudaFreeHost()`

`cudaHostRegister()`, `cudaHostUnregister()`

- Optionally with flags

`cudaHostAllocWriteCombined`

`cudaHostAllocMapped`

`cudaHostAllocPortable`

Optimized for writing,
not cached on CPU

- Copies between pinned host memory and device are automatically performed asynchronously
- Pinned memory is a scarce resource

Memory Mapping

- ▶ Device Memory Mapping
 - Allowing GPU to access portions of host memory directly (i.e., without explicit copy operations)
 - For both reading and writing
 - The memory must be allocated/registered with flag **cudaHostAllocMapped**
 - The context must have **cudaDeviceMapHost** flag (set by **cudaSetDeviceFlags()**)
 - Function **cudaHostGetDevicePointer()** gets host pointer and returns corresponding device pointer

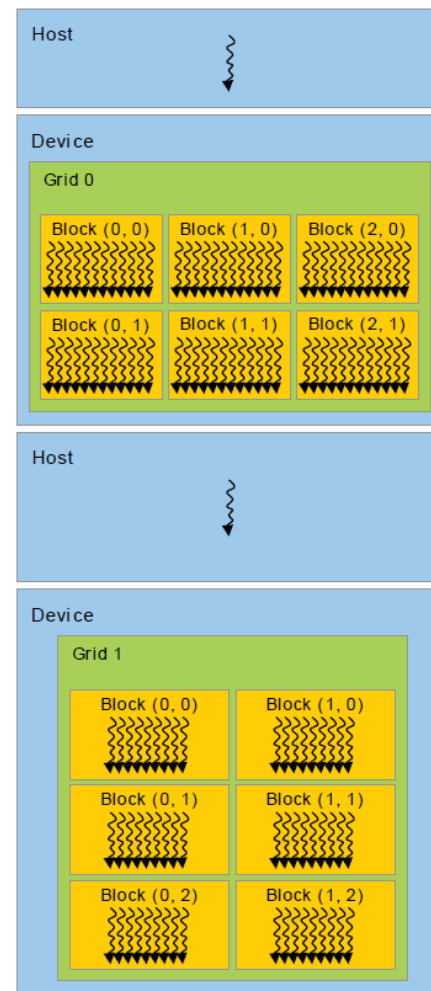
Heterogeneous Programming

▶ GPU

- “Independent” device
- Controlled by host
- Used for “offloading”

▶ Host Code

- Needs to be designed in a way that
 - Utilizes GPU(s) efficiently
 - Utilize CPU while GPU is working
 - CPU and GPU do not wait for each other




Heterogeneous Programming

▶ Bad Example

```
cudaMemcpy(..., HostToDevice);  
Kernel1<<<...>>>(...);  
cudaDeviceSynchronize();  
cudaMemcpy(..., DeviceToHost);  
...  
cudaMemcpy(..., HostToDevice);  
Kernel2<<<...>>>(...);  
cudaDeviceSynchronize();  
cudaMemcpy(..., DeviceToHost);  
...
```

CPU

GPU



Device is doing something useful

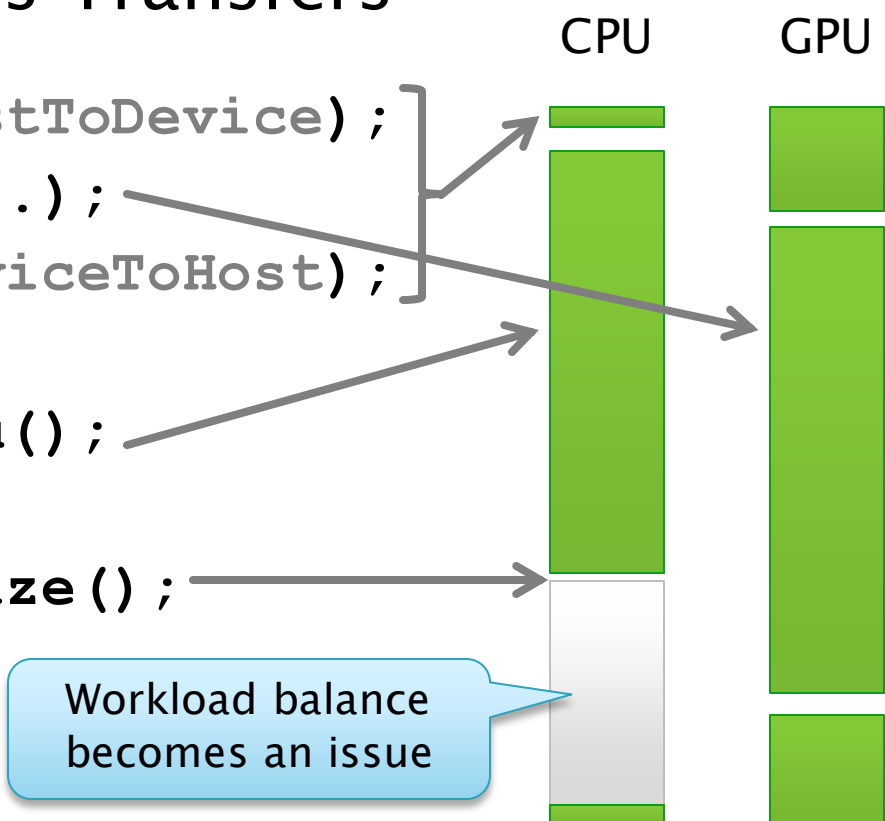
Overlapping Work

- ▶ Overlapping CPU and GPU work
 - Kernels
 - Started asynchronously
 - Can be waited for (`cudaDeviceSynchronize()`)
 - A little more can be done with streams
 - Memory transfers
 - `cudaMemcpy()` is synchronous and blocking
 - Alternatively `cudaMemcpyAsync()` starts the transfer and returns immediately
 - Can be synchronized the same way as the kernel

Overlapping Work

▶ Using Asynchronous Transfers

```
cudaMemcpyAsync (HostToDevice) ;  
Kernel1<<<...>>> (...);  
cudaMemcpyAsync (DeviceToHost) ;  
...  
do_something_on_cpu();  
...  
cudaDeviceSynchronize();
```



Workload balance becomes an issue

Streams

▶ Stream

- In-order GPU command queue
 - Asynchronous GPU operations are registered in queue
 - Kernel execution
 - Memory data transfers
 - Commands in different streams may overlap
 - Provide means for explicit and implicit synchronization
- Default stream (stream 0)
 - Always present, does not have to be created
 - Global synchronization capabilities

Streams

- ▶ Stream Creation

```
cudaStream_t stream;  
cudaStreamCreate(&stream);
```

- ▶ Stream Usage

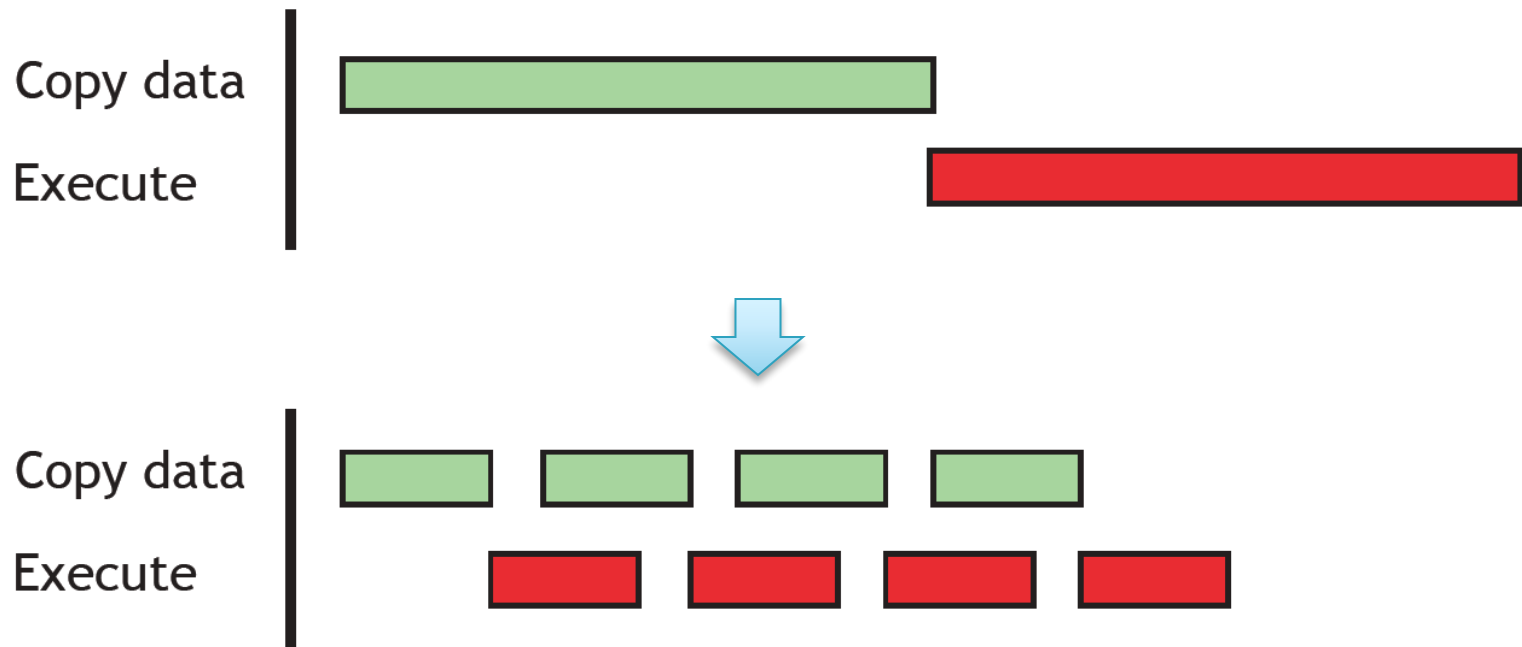
```
cudaMemcpyAsync(dst, src, size, kind, stream);  
kernel<<<grid, block, sharedMem, stream>>>(...);
```

- ▶ Stream Destruction

```
cudaStreamDestroy(stream);
```

Pipelining

- ▶ Making a Good Use of Overlapping
 - Split the work into smaller fragments
 - Create a pipeline effect (load, process, store)



Instruction Set

▶ GPU Instruction Set

- Oriented for rendering and geometry calculations
- Rich set of mathematical functions
 - Many of those are implemented as instructions
 - Separate functions for doubles and floats
 - e.g., `sqrtof(float)` and `sqrt(double)`
- Instruction behavior depends on compiler options
 - `-use_fast_math` – fast but lower precision
 - `-ftz=bool` – flush denormals to zero
 - `-prec-div=bool` – precise float divisions
 - `-prec-sqrt=bool` – precise float sqrts
 - `-fmad=bool` – use mul-add instructions (e.g., FFMA)

Single precision
floats only

Atomic Instructions

▶ Atomic Instructions

- Perform read–modify–write operation of one 32bit or 64bit word in global or shared memory
- Require CC 1.1 or higher (1.2 for 64bit global atomics and 2.0 for 64bit shared mem. atomics)
- Operate on integers, except for **atomicExch()** and **atomicAdd()** which also work on 32bit floats
- Atomic operations on mapped memory are atomic only from the perspective of the device
 - Since they are usually performed on L2 cache

Atomic Instructions

▶ Atomic Instructions Overview

- `atomicAdd(&p, v)`, `atomicSub(&p, v)` – atomically adds or subtracts `v` to/from `p` and return the old value
- `atomicInc(&p, v)`, `atomicDec(&p, v)` – atomic increment/decrement computed modulo `v`
- `atomicMin(&p, v)`, `atomicMax(&p, v)`
- `atomicExch(&p, v)` – atomically swaps a value
- `atomicCAS(&p, v)` – classical compare-and-set
- `atomicAnd(&p, v)`, `atomicOr(&p, v)`, `atomicXor(&p, v)` – atomic bitwise operations

Warp Functions

▶ Voting Instructions

- Intrinsics that allows the whole warp to perform reduction and broadcast in one step
 - Only active threads are participating
- `__all(predicate)`
 - All active threads evaluate predicate
 - Returns non-zero if ALL predicates returned non-zero
- `__any(predicate)`
 - Like `__all`, but the results are combined by logical OR
- `__ballot(predicate)`
 - Return bitmask, where each bit represents the predicate result of the corresponding thread

Warp Functions

▶ Warp Shuffle Instructions

- Fast variable exchange within the warp
- Available for architectures with CC 3.0 or newer
- Intrinsics
 - `__shfl_sync()` – direct copy from given lane
 - `__shfl_up_sync()` – copy with lower relative ID
 - `__shfl_down_sync()` – copy with higher relative ID
 - `__shfl_xor_sync()` – copy from a lane, which ID is computed as XOR of caller ID
 - All functions have optional width parameter, that allows to divide warp into smaller segments

Summary

- ▶ Right amount of threads
 - Saturate SMPs, but avoid registry spilling
- ▶ SIMT
 - Avoid warp divergence
 - Synchronization within a block is cheap
- ▶ Memory
 - Host–device transfer overlapping
 - Global memory coalesced transactions
 - Shared memory banking

Discussion

